SAMPLE 4 OF 4

ServiceFlow / TechDell -- Agentic AI Workflow Orchestration Safety Card

Platform Ref	AW-3 Series -- ServiceFlow Now Platform AI, integrated with TechDell Partner Resource Portal . Documentation Cycle: Q1-Q2 2024
Card Version	v1.1 . Joint Customer Safety Disclosure . Pilot deployment to 8 enterprise customers
Author	E. Hannan -- Documentation Lead & Research Operations Coordinator; HAX outside-in methodology applied to user testing
Source Team	ServiceFlow AI Platform Engineering . TechDell Partner Enablement . Joint Customer IT Security . External Policy Vulnerability Testing Partner
Audience	Enterprise IT Architects . Partner Program Leadership . AI Governance Committees . Compliance Counsel
Classification	Joint Partner Release -- NDA Protected . Names adapted per disclosure protocol
Related Artifacts	ServiceFlow Now Platform AI Governance Brief . TechDell Partner Portal Security Assessment . Agentic AI Policy Framework (Draft)
Software	ServiceFlow Now Platform Washington DC GA (Q3 2023) . Now Assist (GenAI, agentic ITSM) . TechDell Partner Portal v4.2 (Q4 2023)

1. Executive Summary

The AW-3 Series describes an agentic AI orchestration layer deployed across the ServiceFlow Now Platform and the TechDell Partner Resource Portal -- enabling automated IT service management, partner deal progression, and cross-platform workflow execution with reduced human-in-the-loop requirements at each step. This documentation cycle began with the pilot deployment of ServiceFlow Washington DC (GA Q3 2023) and TechDell Partner Portal v4.2 (GA Q4 2023) to eight enterprise customers.

This card documents a deployment at the frontier of enterprise agentic AI: systems where an AI model takes multi-step actions across integrated platforms, with each action potentially triggering downstream consequences that are difficult to reverse. The safety architecture required for this class of deployment is meaningfully higher than for single-turn AI assistants. We write this card to that higher standard because the consequences of overstating safety properties for agentic systems are qualitatively different -- irreversible actions cause harm that post-hoc remediation cannot fully address.

2. Agentic System Capabilities

2.1 Capability Scope -- Production and Pilot

Evaluation Metric	Result	Documentation Note
Autonomous IT ticket triage and routing	Active in production	85.3% accuracy at routing to correct tier-1 resolution path; human override rate 4.1% of routed tickets; baseline measured over 90 days
Multi-step change management orchestration	Active in production	Requires explicit human approval at Change Advisory Board stage; agent cannot advance past this gate autonomously
Partner deal registration (ServiceFlow to TechDell portal)	Active in production	Cross-system write access; permission scope bounded per Section 3; reversibility window 72 hours
Automated compute resource provisioning	Pilot only	Sandboxed environment; no production system write access during pilot phase
Natural language policy query and interpretation	Active in production	Advisory output only; no autonomous policy decisions; all outputs logged and auditable
Cross-platform agent chaining (3+ sequential steps)	Pilot only	Highest risk class; adversarial testing of multi-step chains not yet complete -- see Section 4.2

2.2 The Agentic Safety Threshold

Agentic AI systems require a different safety analysis than single-turn assistants because the action surface is cumulative: each step an agent takes narrows the space of possible reversals and expands the potential for downstream harm. The capability table above reflects deliberate deployment scope decisions, not technical limitations. We chose to keep automated compute provisioning and cross-platform agent chaining in pilot status because the adversarial testing required to justify production deployment had not been completed at the time this card was produced.

3. Safety Architecture -- Bounded Agency

- Minimal footprint: the agent requests only the permissions required for the current task; permissions are not retained between sessions and are not elevated based on inferred task context
- Explicit scope ceiling: write access is limited to ServiceFlow ITSM records and TechDell deal registration tables; no access to financial systems, HR systems, or customer PII databases in any current deployment configuration
- Reversibility first: all agentic actions that modify records produce a prior-state snapshot; rollback is available for 72 hours post-action; rollback availability is surfaced in the action confirmation UI

- Human confirmation gates: Change Advisory Board approval, partner tier upgrades, and compute provisioning requests require explicit human confirmation regardless of agent confidence score
- Anomaly detection overlay: action sequences exceeding volume thresholds, targeting atypical resource classes, or occurring outside established usage windows trigger a pause-and-notify protocol before further execution
- Policy Vulnerability Testing (PVT): external partner conducted structured adversarial testing across 140 scenarios in 7 risk categories prior to pilot launch; findings documented in Section 4

4. Risk Surface & Honest Disclosures

4.1 Policy Vulnerability Testing Results -- 140 Scenarios, 7 Risk Categories

Evaluation Metric	Result	Documentation Note
Scenarios tested	140 total	Across 7 risk categories: permission escalation, data exfiltration, change management bypass, cross-system injection, session persistence, rollback evasion, social engineering of human gates
Critical findings	0	No scenario produced an irreversible high-impact action during the testing period
High findings	2	Both remediated before pilot launch -- see details below
Medium findings	4	3 remediated; 1 deferred with compensating control and documented rationale
Low findings	6	All documented; monitoring in place; none escalated to medium during 90-day pilot

High Finding 1 -- Permission Scope Escalation via Ticket Notes: The agent could be prompted through free-text input in the ticket notes field to request permissions outside its defined scope ceiling. This is a prompt injection vulnerability in an agentic context -- categorically more serious than in a single-turn context because the elevated permission, once granted, persists across the action sequence. Remediated via strict input sanitization and scope validation at the permission-request layer before any API call is made.

High Finding 2 -- Silent Permission Escalation on Retry: In a specific multi-step sequence where an action failed due to insufficient permissions, the agent would retry with elevated permissions without surfacing a notification to the user. This violates the minimal footprint principle and removes the human from the escalation decision. Remediated by adding mandatory user notification on all permission escalation attempts, with explicit user confirmation required before retry.

4.2 Cross-Platform Agent Chaining (Pilot Only -- Active Risk)

Scenarios involving three or more sequential actions across both platforms remain in pilot status because the risk surface compounds at each step in ways that single-step or two-step adversarial testing does not adequately capture. We recommended -- and the ServiceFlow and TechDell product teams accepted -- that this capability class should not advance to general availability until adversarial testing specifically targeting multi-step chain sequences is complete. This recommendation was made against commercial pressure to accelerate the full GA timeline. The documentation function is an accountability function; this recommendation stands because the testing has not been done.

5. Disability Access & Inclusive Design Audit

Agentic AI systems that take actions on behalf of users create new accessibility stakes that single-turn assistants do not. A user who relies on screen reader technology or alternative input devices must be able to review, confirm, and override agentic actions with the same confidence and with the same information as a user operating without assistive technology. If the confirmation dialog is not screen-reader accessible, that user cannot meaningfully consent to the action being taken in their name. We audited the AW-3 deployment against this standard and identified two gaps before pilot launch:

- The confirmation dialog for agentic actions was not screen-reader accessible in the ServiceFlow mobile interface: the dialog rendered as a modal with no ARIA role assignment, no announced focus management, and no accessible dismiss path
- The action history log did not support keyboard-only navigation: users could not tab through the action history to review or initiate rollback without mouse interaction

Both gaps were remediated before pilot launch. This audit was initiated by the documentation lead, not by the product team's standard QA process, based on prior accessibility and disability research work. Equitable access to the oversight layer of an agentic system is not a compliance requirement -- it is a precondition for meaningful human control.

6. On the Responsibility of Documenting Agentic Systems

I want to say something directly, in my own voice, about what it means to write safety documentation for agentic AI systems. The research I have produced on sycophancy, prompt injection, and ESG adversarial robustness gives me a specific lens on this work: documentation that overstates a system's safety properties does not merely create legal liability -- it causes real harm when the system fails in ways the documentation said it would not. In the context of systems that take irreversible actions, this is not a theoretical concern.

The two high findings in Section 4.1 were surfaced during documentation review, not during engineering review. The documentation process -- asking what happens when a bad actor controls the input, reading the system as an adversary would, refusing to accept confident assertions about safety without evidence -- is itself a safety function. I did not learn this from a framework. I learned it from 18 years of producing accountability documentation for complex enterprise systems that real organizations

depend on, and from the consistent experience that the most important findings surface when you refuse to accept the comfortable answer.